Hadoop+Spark+ iObjects for Spark

部署流程 v0.2

一、 准备

1.1 软件版本

- Ubuntu 14.04.4
- JDK 1.7.0
- Scala 2.10.6
- Hadoop 2.6.5
- Spark 1.6.3
- SuperMap iObjects Java 8.1.1
- SuperMap iObjects for Spark 8.1.1_60034

1.2 SSH 设置

修改/etc/ssh/sshd_config 文件，将以下三项开启 yes 状态

```
PermitRootLogin yes
PermitEmptyPasswords yes
PasswordAuthentication yes
```

重启 ssh 服务

```
service ssh restart
```

这样 root 用户可直接登陆，以及为后续 ssh 无密码登录做准备。

1.3 绑定 IP 和修改计算机名

1.3.1 修改/etc/hosts,添加 IP 绑定注释 127.0.1.1 绑定（不注释会影响 hadoop 集群）

```
127.0.0.1       localhost
#127.0.1.1       sparkalone

192.168.15.29   sparkalone-master
```

1.3.2 修改/etc/hostname,为绑定计算机名。(计算机名和上面 hosts 绑定名必须一致)

1.4 SSH 无密码登陆

1、cd ~/.ssh　　#进到当前用户的隐藏目录（.ssh）

2、ssh-keygen -t rsa　　#用 rsa 生成密钥，一路回车

3、cp id_rsa.pub authorized_keys　　#把公钥复制一份，并改名为 authorized_keys，这步执行完后，在当前机器执行 ssh localhost 可以无密码登录本机了

*初次执行会提示确认，输入 yes 和登陆密码，之后就没提示了。

*后续如果做集群，可以 scp authorized_keys root@第二台机器名：/root/.ssh　（把重命名后的公钥通过 ssh 提供的远程复制文件复制到从机）这样子节点启动也不需要再次输入密码。

1.5 JDK 安装

下载：jdk-7u80-linux-x64.tar.gz 包，放到/opt 下解压

1.5.1 将 JDK 环境变量配置到/etc/profile 中

```
export JAVA_HOME=/opt/jdk1.7.0_80
export JRE_HOME=/opt/jdk1.7.0_80/jre
export CLASSPATH=$JAVA_HOME/lib:$JRE_HOME/lib
export PATH=$JAVA_HOME/bin: $PATH
```

1.5.2 检查 JDK 是否配置好

```
root@sparkalone-master:~# java -version
java version "1.7.0_80"
Java(TM) SE Runtime Environment (build 1.7.0_80-b15)
Java HotSpot(TM) 64-Bit Server VM (build 24.80-b11, mixed mode)
```

二、 Hadoop 部署

2.1 Hadoop 安装

1. 下载 hadoop2.6.5（hadoop-2.6.5.tar.gz）
2. 解压 tar -zxvf hadoop-2.6.5.tar.gz，并在主目录下创建 tmp、dfs、dfs/name、dfs/node、dfs/data

```
root@sparkalone-master:/opt/hadoop-2.6.5# mkdir tmp
root@sparkalone-master:/opt/hadoop-2.6.5# mkdir dfs
```

```
root@sparkalone-master:/opt/hadoop-2.6.5# mkdir dfs/name
root@sparkalone-master:/opt/hadoop-2.6.5# mkdir dfs/node
root@sparkalone-master:/opt/hadoop-2.6.5# mkdir dfs/data
```

2.2 Hadoop 配置

以下操作都在 hadoop-2.6.5/etc/hadoop 下进行

2.2.1 编辑 hadoop-env.sh 文件，修改 JAVA_HOME 配置项为 JDK 安装目录

```
export JAVA_HOME=/opt/jdk1.7.0_80
```

2.2.2 编辑 core-site.xml 文件，添加以下内容，

其中 **sparkalone-master** 为计算机名，**/opt/hadoop-2.6.5/tmp** 为手动创建的目录

```
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://sparkalone-master:9000</value>
  </property>
  <property>
    <name>io.file.buffer.size</name>
    <value>131072</value>
  </property>
  <property>
    <name>hadoop.tmp.dir</name>
    <value>file:/opt/hadoop-2.6.5/tmp</value>
    <description>Abasefor other temporary directories.</description>
  </property>
  <property>
    <name>hadoop.proxyuser.spark.hosts</name>
    <value>*</value>
  </property>
<property>
    <name>hadoop.proxyuser.spark.groups</name>
    <value>*</value>
  </property>
</configuration>
```

2.2.3 编辑 hdfs-site.xml 文件，添加以下内容

其中 **sparkalone-master** 为计算机名，
file:/opt/hadoop-2.6.5/dfs/name 和 **file:/opt/hadoop-2.6.5/dfs/data** 为手动创建目录

```
<configuration>
<property>
   <name>dfs.namenode.secondary.http-address</name>
   <value>sparkalone-master:9001</value>
 </property>
   <property>
     <name>dfs.namenode.name.dir</name>
     <value>file:/opt/hadoop-2.6.5/dfs/name</value>
 </property>
 <property>
   <name>dfs.datanode.data.dir</name>
   <value>file:/opt/hadoop-2.6.5/dfs/data</value>
   </property>
 <property>
   <name>dfs.replication</name>
   <value>3</value>
 </property>
 <property>
   <name>dfs.webhdfs.enabled</name>
   <value>true</value>
 </property>
</configuration>
```

复制 mapred-site.xml.template 并重命名为 mapred-site.xml

```
cp mapred-site.xml.template mapred-site.xml
```

2.2.4 编辑 mapred-site.xml 文件，添加以下内容

其中 **sparkalone-master** 为计算机名

```
<configuration>
<property>
     <name>mapreduce.framework.name</name>
     <value>yarn</value>
 </property>
 <property>
   <name>mapreduce.jobhistory.address</name>
   <value>sparkalone-master:10020</value>
 </property>
 <property>
   <name>mapreduce.jobhistory.webapp.address</name>
   <value>sparkalone-master:19888</value>
 </property>
</configuration>
```

2.2.5 编辑 yarn-site.xml 文件，添加以下内容

其中 **sparkalone-master** 为计算机名

```
<configuration>
<!-- Site specific YARN configuration properties -->
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
  <property>
    <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
    <value>org.apache.hadoop.mapred.ShuffleHandler</value>
  </property>
  <property>
    <name>yarn.resourcemanager.address</name>
    <value>sparkalone-master:8032</value>
  </property>
  <property>
    <name>yarn.resourcemanager.scheduler.address</name>
    <value>sparkalone-master:8030</value>
  </property>
  <property>
    <name>yarn.resourcemanager.resource-tracker.address</name>
    <value>sparkalone-master:8035</value>
  </property>
  <property>
    <name>yarn.resourcemanager.admin.address</name>
    <value>sparkalone-master:8033</value>
  </property>
  <property>
    <name>yarn.resourcemanager.webapp.address</name>
    <value>sparkalone-master:8088</value>
  </property>
</configuration>
```

2.2.6 修改 slaves 文件，添加本机作为集群节点

```
root@sparkalone-master:/opt/hadoop-2.6.5/etc/hadoop# cat slaves
sparkalone-master
```

注意：

如果 hadoop 配置集群，可以将配置文件 etc/hadoop 下内容同步到其他机器上，并修改 slaves 文件

```
scp -r hadoop root@另一台机器名:/opt /hadoop-2.6.5/etc
```

2.3 Hadoop 启动

1. 格式化一个新的文件系统，进入到 hadoop-2.6.5/bin 下执行：

```
./hadoop namenode –format
```

2. **启动 hadoop**，进入到 hadoop-2.6.5/sbin 下执行：

```
./start-all.sh
```

看到如下内容说明启动成功

```
root@sparkalone-master:/opt/hadoop-2.6.5/sbin# ./start-all.sh
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh
Starting namenodes on [sparkalone-master]
sparkalone-master: starting namenode, logging to /opt/hadoop-2.6.5/logs/hadoop-root-namenode-sparkalone-master.out
sparkalone-master: starting datanode, logging to /opt/hadoop-2.6.5/logs/hadoop-root-datanode-sparkalone-master.out
Starting secondary namenodes [sparkalone-master]
sparkalone-master: starting secondarynamenode, logging to /opt/hadoop-2.6.5/logs/hadoop-root-secondarynamenode-sparkalone-master.out
starting yarn daemons
starting resourcemanager, logging to /opt/hadoop-2.6.5/logs/yarn-root-resourcemanager-sparkalone-master.out
sparkalone-master: starting nodemanager, logging to /opt/hadoop-2.6.5/logs/yarn-root-nodemanager-sparkalone-master.out
```

2.4 Hadoop 集群检查

检查 hadoop 集群，进入 hadoop-2.6.5/bin 下执行

```
./hdfs dfsadmin -report
```

```
root@sparkalone-master:/opt/hadoop-2.6.5/bin# ./hdfs dfsadmin -report
Configured Capacity: 50596724736 (47.12 GB)
Present Capacity: 42564304896 (39.64 GB)
DFS Remaining: 42490081280 (39.57 GB)
DFS Used: 74223616 (70.79 MB)
DFS Used%: 0.17%
Under replicated blocks: 2
Blocks with corrupt replicas: 0
Missing blocks: 0

-------------------------------------------------
Live datanodes (1):

Name: 192.168.15.29:50010 (sparkalone-master)
Hostname: sparkalone-master
Decommission Status : Normal
Configured Capacity: 50596724736 (47.12 GB)
DFS Used: 74223616 (70.79 MB)
Non DFS Used: 8032419840 (7.48 GB)
DFS Remaining: 42490081280 (39.57 GB)
DFS Used%: 0.15%
DFS Remaining%: 83.98%
Configured Cache Capacity: 0 (0 B)
Cache Used: 0 (0 B)
Cache Remaining: 0 (0 B)
Cache Used%: 100.00%
Cache Remaining%: 0.00%
Xceivers: 1
Last contact: Wed Jan 18 14:13:27 CST 2017
```

或者访问 http://192.168.15.29:8088/cluster/nodes

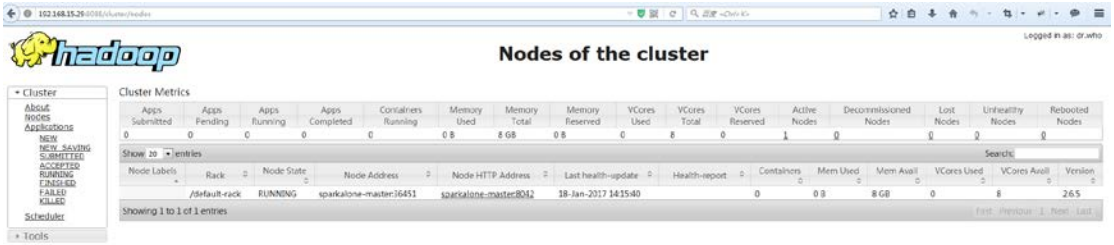

三、 Spark 部署

3.1 Spark 安装

1. 下载 scala-2.10.6.tgz 和 spark-1.6.3-bin-hadoop2.6.tgz，放到 opt 下解压。
2. 将 scala 和 spark 环境变量配置到/etc/profile 中

```
export SCALA_HOME=/opt/scala-2.10.6
export SPARK_HOME=/opt/spark-1.6.3-bin-hadoop2.6
export PATH=$SCALA_HOME/bin:$JAVA_HOME/bin:$SPARK_HOME/bin:$PATH
```

3.2 Spark 配置

1. 进入 spark-1.6.3-bin-hadoop2.6/conf

复制 spark-env.sh.template 并重命名为 spark-env.sh

```
cp spark-env.sh.template spark-env.sh
```

编辑 spark-env.sh 文件，添加以下内容

```
export JAVA_HOME=/opt/jdk1.7.0_80
export SPARK_MASTER_IP=192.168.15.29
export SPARK_WORKER_MEMORY=4g
export SCALA_HOME=/opt/scala-2.10.6
export HADOOP_HOME=/opt/hadoop-2.6.5/
export HADOOP_CONF_DIR=/opt/hadoop-2.6.5/etc/hadoop
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$SUPERMAP_OBJ
export SPARK_CLASSPATH=$SPARK_CLASSPATH
```

2.编辑 slaves 文件，增加当前计算机名

```
root@sparkalone-master:/opt/spark-1.6.3-bin-hadoop2.6/conf# cat slaves
#
# Licensed to the Apache Software Foundation (ASF) under one or more
# contributor license agreements.  See the NOTICE file distributed with
# this work for additional information regarding copyright ownership.
# The ASF licenses this file to You under the Apache License, Version 2.0
# (the "License"); you may not use this file except in compliance with
# the License.  You may obtain a copy of the License at
#
#    http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.
#

# A Spark Worker will be started on each of the machines listed below.
sparkalone-master
```

3.3 Spark 启动

启动 spark，进入 spark-1.6.3-bin-hadoop2.6/sbin 下执行

```
./start-all.sh
```

3.4 Spark 集群检查

访问 http://192.168.15.29:8080/

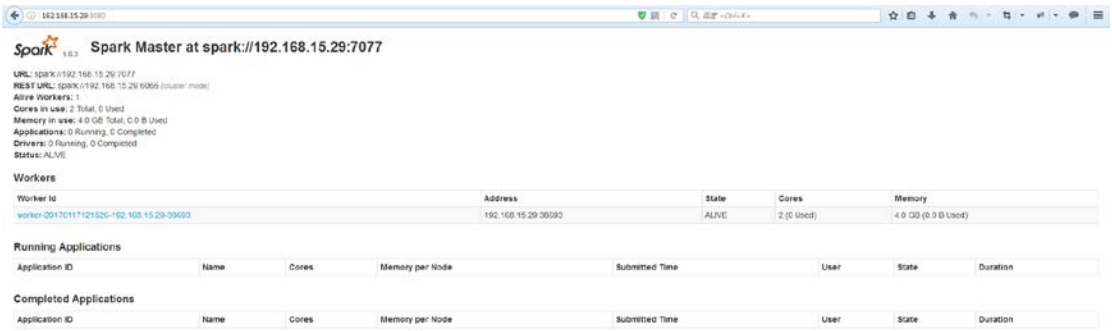

注意：如果将来配置 spark 集群，只需要保证子节点内容和主节点内容一致即可。

四、iObjects for Spark 部署

4.1 iObjects for Spark 安装

1. 下载 supermap_iobjects_for_spark_811_60034.tar.gz 和相应的 iObjects Java 811 组件
2. 放到 opt 下解压

（例如 for spark 中 lib 放到/opt/iobjects_spark，iObjects Java 放到/opt/iobjects_new）

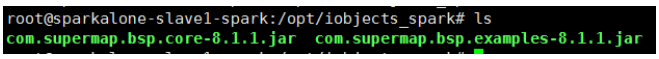

4.2 iObjects for Spark 配置

1. 将 for Spark 和 iObjects Java 环境变量配置到/etc/profile 中,结合之前 hadoop 和 spark 配置。总配置如下，其中红色为新增加配置。

```
export JAVA_HOME=/opt/jdk1.7.0_80
export SCALA_HOME=/opt/scala-2.10.6
export SPARK_HOME=/opt/spark-1.6.3-bin-hadoop2.6
export PATH=$SCALA_HOME/bin:$JAVA_HOME/bin:$SPARK_HOME/bin:$PATH
export SUPERMAP_OBJ=/opt/iobjects_new
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$SUPERMAP_OBJ
export SPARK_CLASSPATH=$SPARK_CLASSPATH:/opt/iobjects_spark
export JRE_HOME=/opt/jdk1.7.0_80/jre
export CLASSPATH=$JAVA_HOME/lib:$JRE_HOME/lib
```

2. 进入 spark-1.6.3-bin-hadoop2.6/conf

编辑 spark-env.sh 文件，总配置如下，其中红色为新增加配置。

```
export JAVA_HOME=/opt/jdk1.7.0_80
export SPARK_MASTER_IP=192.168.15.29
export SPARK_WORKER_MEMORY=4g
export SCALA_HOME=/opt/scala-2.10.6
export HADOOP_HOME=/opt/hadoop-2.6.5/
export HADOOP_CONF_DIR=/opt/hadoop-2.6.5/etc/hadoop
export SUPERMAP_OBJ=/opt/iobjects_new
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$SUPERMAP_OBJ
export SPARK_CLASSPATH=$SPARK_CLASSPATH:/opt/iobjects_spark
```

4.3 实例程序验证

我们使用 for Spark 示例来验证是否安装成功。

1. 将产品包中的示例数据放到中/opt/data 文件夹中，并创建/opt/resultData 空文件夹

```
root@sparkalone-master:/opt/data# ls
taxi.json  taxi.meta
```

```
root@sparkalone-master:/opt/resultData# ls
root@sparkalone-master:/opt/resultData#
```

2. 将示例数据导入到 hdfs 中。

   启动 hadoop，在 hadoop-2.6.5/bin 中执行

```
./hadoop fs -mkdir /input       #创建/input 目录
./hdfs dfs -put /opt/data/taxi.json /input/     #将 taxi.json 导入到/input 目录中
./hdfs dfs -put /opt/data/taxi.meta /input/    #将 taxi. meta 导入到/input 目录中
```

   导入完成后，可以使用如下命名查看

```
./hadoop fs -ls /input
```

```
root@sparkalone-master:/opt/hadoop-2.6.5/bin# ./hadoop fs -ls /input
Found 2 items
-rw-r--r--   3 root supergroup   73595399 2017-01-17 14:52 /input/taxi.json
-rw-r--r--   3 root supergroup       2451 2017-01-17 15:07 /input/taxi.meta
```

3. 启动 Spark，切换到 spark-1.6.3-bin-hadoop2.6/bin 下执行（执行几何对象创建缓冲区）

```
./spark-submit --class com.supermap.bsp.examples.GeometryBufferTest --master spark://192.168.15.29:7077 --conf spark.default.parallelism=3 --jars /opt/iobjects_spark/com.supermap.bsp.core-8.1.1.jar /opt/iobjects_spark/com.supermap.bsp.examples-8.1.1.jar --input hdfs://sparkalone-master:9000/input/taxi.json --distance "10 meter" --output /opt/resultData/BufferResult.json
```

简单解释下命令

--class　　表示主类名称，含包名，本例子指的是需要执行的类

--master　　Spark 集群总入口

--conf　　spark.default.parallelism 值跟集群的总核数有关，建议是总核数的两到三倍

--jars　　依赖的第三方 jar 包

--input　　简单理解为操作数据来源

--distance　　参数，这里只范例中缓冲距离

--output　　结果输出到哪里

注意：如果集群模式运行不成功，
将-master spark://192.168.15.29:7077 修改为：-master local[4]本地运行
如果 local 模式运行还是报错，请参照常见问题 5.8

执行开始

```
root@sparkalone-master:/opt/spark-1.6.3-bin-hadoop2.6/bin# ./spark-submit --class com.supermap.bsp.examples.GeometryBufferTest --master spark://192.168.15.29:7077 --conf spark.default.parallelism=3 --jars /opt/iobjects_spark/com.supermap.bsp.core-8.1.1.jar /opt/iobjects_spark/com.supermap.bsp.examples-8.1.1.jar --input hdfs://sparkalone-master:9000/input/taxi.json --distance "10 meter" --output /opt/resultData/BufferResult.json
17/01/18 13:32:06 INFO spark.SparkContext: Running Spark version 1.6.3
17/01/18 13:32:06 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
17/01/18 13:32:07 WARN spark.SparkConf:
SPARK_CLASSPATH was detected (set to ':/opt/iobjects_spark:/opt/iobjects_spark:/opt/iobjects_spark').
This is deprecated in Spark 1.0+.

Please instead use:
 - ./spark-submit with --driver-class-path to augment the driver classpath
 - spark.executor.extraClassPath to augment the executor classpath

17/01/18 13:32:07 WARN spark.SparkConf: Setting 'spark.executor.extraClassPath' to ':/opt/iobjects_spark:/opt/iobjects_spark:/opt/iobjects_spark' as a work-around.
17/01/18 13:32:07 WARN spark.SparkConf: Setting 'spark.driver.extraClassPath' to ':/opt/iobjects_spark:/opt/iobjects_spark:/opt/iobjects_spark' as a work-around.
17/01/18 13:32:07 INFO spark.SecurityManager: Changing view acls to: root
17/01/18 13:32:07 INFO spark.SecurityManager: Changing modify acls to: root
17/01/18 13:32:07 INFO spark.SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view permissions: Set(root); users with modify permissions: Set(root)
17/01/18 13:32:08 INFO util.Utils: Successfully started service 'sparkDriver' on port 34177.
17/01/18 13:32:08 INFO slf4j.Slf4jLogger: Slf4jLogger started
17/01/18 13:32:08 INFO Remoting: Starting remoting
17/01/18 13:32:09 INFO Remoting: Remoting started; listening on addresses :[akka.tcp://sparkDriverActorSystem@192.168.15.29:39156]
17/01/18 13:32:09 INFO util.Utils: Successfully started service 'sparkDriverActorSystem' on port 39156.
17/01/18 13:32:09 INFO spark.SparkEnv: Registering MapOutputTracker
```

执行过程中，可以访问 http://192.168.15.29:4040/jobs/ 查看执行情况

执行完成后，访问/opt/resultData 可以查看完成的内容

```
root@sparkalone-master:/opt/resultData# ls
BufferResult.json  BufferResult.meta
```

五、常见问题

5.1 启动 hadoop 来回让输入密码

答：没设置免 ssh 密码登陆， 截图如下，解决办法请参考 1.2 和 1.4

```
root@sparkalone-master:/opt/hadoop-2.6.5/sbin# ./start-all.sh
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh
Starting namenodes on [sparkalone-master]
The authenticity of host 'sparkalone-master (192.168.15.29)' can't be established.
ECDSA key fingerprint is 09:21:b7:bf:f9:03:e4:d2:d9:04:5b:d8:92:9a:7e:71.
Are you sure you want to continue connecting (yes/no)? yes
sparkalone-master: Warning: Permanently added 'sparkalone-master,192.168.15.29' (ECDSA) to the list of known hosts.
root@sparkalone-master's password:
sparkalone-master: starting namenode, logging to /opt/hadoop-2.6.5/logs/hadoop-root-namenode-sparkalone-master.out
parkalone-master: ssh: Could not resolve hostname parkalone-master: Name or service not known
Starting secondary namenodes [sparkalone-master]
root@sparkalone-master's password:
sparkalone-master: starting secondarynamenode, logging to /opt/hadoop-2.6.5/logs/hadoop-root-secondarynamenode-sparkalone-master.out
starting yarn daemons
starting resourcemanager, logging to /opt/hadoop-2.6.5/logs/yarn-root-resourcemanager-sparkalone-master.out
parkalone-master: ssh: Could not resolve hostname parkalone-master: Name or service not known
root@sparkalone-master:/opt/hadoop-2.6.5/sbin#
```

5.2 如何判断 hadoop 集群是否启动成功

答：可以，进入 hadoop-2.6.5/bin 下执行

```
./hdfs dfsadmin -report
```

下图就是没启动成功，具体原因可以查看 hadoop-2.6.5/logs 日志

```
root@sparkalone-master:/opt/hadoop-2.6.5/bin# ./hdfs dfsadmin -report
Configured Capacity: 0 (0 B)
Present Capacity: 0 (0 B)
DFS Remaining: 0 (0 B)
DFS Used: 0 (0 B)
DFS Used%: NaN%
Under replicated blocks: 0
Blocks with corrupt replicas: 0
Missing blocks: 0
```

5.3 还有什么命令可以查看集群是否启动成功

答：输入 jps，根据启动的进程，查看是否启动成功

```
root@sparkalone-master:/opt/spark-1.6.3-bin-hadoop2.6/sbin# jps
3081 SecondaryNameNode
3353 NodeManager
3221 ResourceManager
3922 Master
2740 NameNode
4109 Jps
2894 DataNode
4026 Worker
```

其中

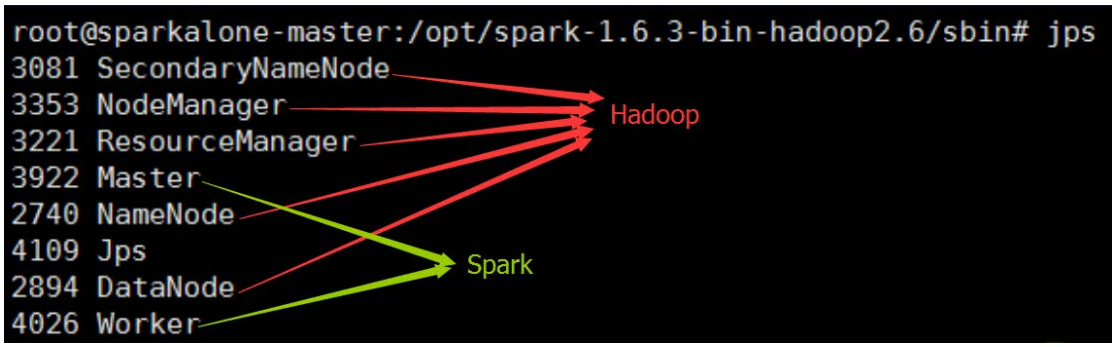

5.4 如何解压 zip 文件

答：需要安装 unzip 命令，安装完成后使用 unzip+文件全名方式解压

```
root@sparkalone-master:/opt/iobjects# apt-get install unzip
Reading package lists... Done
Building dependency tree
Reading state information... Done
Suggested packages:
  zip
The following NEW packages will be installed:
  unzip
0 upgraded, 1 newly installed, 0 to remove and 114 not upgraded.
Need to get 157 kB of archives.
After this operation, 395 kB of additional disk space will be used.
Get:1 http://cn.archive.ubuntu.com/ubuntu/ trusty-updates/main unzip amd64 6.0-9ubuntu1.5 [157 kB]
Fetched 157 kB in 0s (387 kB/s)
Selecting previously unselected package unzip.
(Reading database ... 57572 files and directories currently installed.)
Preparing to unpack .../unzip_6.0-9ubuntu1.5_amd64.deb ...
Unpacking unzip (6.0-9ubuntu1.5) ...
Processing triggers for man-db (2.6.7.1-1ubuntu1) ...
Processing triggers for mime-support (3.54ubuntu1.1) ...
Setting up unzip (6.0-9ubuntu1.5) ...
```

5.5 如何有界面方式查看 hdfs 存储大小

答：访问 http://192.168.15.29:50070/dfshealth.html#tab-datanode

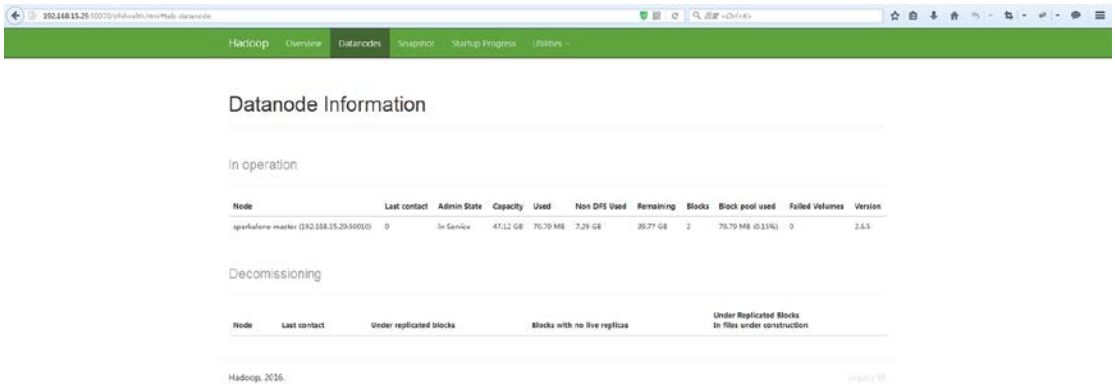

5.6 Spark 启动报错如何查看

答：例如如下 Spark 启动报错，查看 sparkalone-slave1-spark 机器 Spark 下 logs 日志即可

```
root@sparkalone-master:/opt/spark-1.6.3-bin-hadoop2.6/sbin# ./start-all.sh
starting org.apache.spark.deploy.master.Master, logging to /opt/spark-1.6.3-bin-hadoop2.6/logs/spark-root-org.apache.spark.deploy.master.Master-1-sparkalone-master.out
sparkalone-master: starting org.apache.spark.deploy.worker.Worker, logging to /opt/spark-1.6.3-bin-hadoop2.6/logs/spark-root-org.apache.spark.deploy.worker.Worker-1-sparkalone-master.out
sparkalone-slave1-spark: starting org.apache.spark.deploy.worker.Worker, logging to /opt/spark-1.6.3-bin-hadoop2.6/logs/spark-root-org.apache.spark.deploy.worker.Worker-1-parkalone-slave1-spark.out
sparkalone-slave1-spark: failed to launch org.apache.spark.deploy.worker.Worker:
sparkalone-slave1-spark: full log in /opt/spark-1.6.3-bin-hadoop2.6/logs/spark-root-org.apache.spark.deploy.worker.Worker-1-parkalone-slave1-spark.out
```

5.7 Spark 集群启动成功和集群执行过程如何查看

答：访问 http://192.168.15.29:8080/

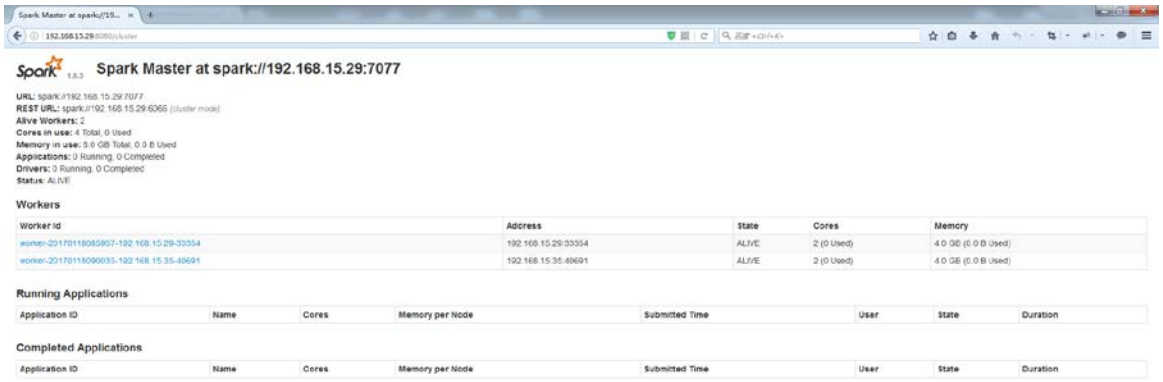

执行过程截图：

访问 http://192.168.15.29:4040/jobs/

5.8 iObjects for Spark 程序使用 local 模式运行不成功，报错

java.lang.NoClassDefFoundError: Could not initialize class java.awt.Toolkit
java.lang.NoClassDefFoundError: Could not initialize class com.supermap.geotools.Convertor

```
17/02/07 23:03:31 ERROR executor.Executor: Exception in task 1.0 in stage 0.0 (TID 1)
java.lang.NoClassDefFoundError: Could not initialize class java.awt.Toolkit
        at java.awt.Component.<clinit>(Component.java:595)
        at com.supermap.data.Environment.LoadWrapJ(Environment.java:375)
        at com.supermap.data.Environment.<clinit>(Environment.java:30)
        at com.supermap.geotools.LocalLogger.initConfig(LocalLogger.java:72)
        at com.supermap.geotools.LocalLogger.<clinit>(LocalLogger.java:14)
        at com.supermap.geotools.Convertor.<clinit>(Convertor.java:53)
        at com.supermap.bsp.rdd.BufferFeatureRDD.compute(BufferFeatureRDD.scala:27)
        at org.apache.spark.rdd.RDD.computeOrReadCheckpoint(RDD.scala:306)
        at org.apache.spark.rdd.RDD.iterator(RDD.scala:270)
        at org.apache.spark.scheduler.ResultTask.runTask(ResultTask.scala:66)
        at org.apache.spark.scheduler.Task.run(Task.scala:89)
        at org.apache.spark.executor.Executor$TaskRunner.run(Executor.scala:227)
        at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:114
5)
        at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:61
5)
        at java.lang.Thread.run(Thread.java:745)
```

```
17/02/07 22:55:43 INFO scheduler.TaskSetManager: Starting task 1.0 in stage 0.0 (TID 1,
Ubuntu, partition 1,NODE_LOCAL, 2286 bytes)
17/02/07 22:55:46 INFO storage.BlockManagerInfo: Added broadcast_3_piece0 in memory on U
buntu:44329 (size: 4.0 KB, free: 511.5 MB)
17/02/07 22:55:51 INFO storage.BlockManagerInfo: Added broadcast_2_piece0 in memory on U
buntu:44329 (size: 339.0 B, free: 511.5 MB)
17/02/07 22:55:53 WARN scheduler.TaskSetManager: Lost task 1.0 in stage 0.0 (TID 1, Ubun
tu): java.lang.UnsatisfiedLinkError: com.supermap.data.EnvironmentNative.jni_GetBasePath
()Ljava/lang/String;
        at com.supermap.data.EnvironmentNative.jni_GetBasePath(Native Method)
        at com.supermap.data.Environment.getUGOBasePath(Environment.java:531)
        at com.supermap.geotools.LocalLogger.initConfig(LocalLogger.java:72)
        at com.supermap.geotools.LocalLogger.<clinit>(LocalLogger.java:14)
        at com.supermap.geotools.Convertor.<clinit>(Convertor.java:53)
        at com.supermap.bsp.rdd.BufferFeatureRDD.compute(BufferFeatureRDD.scala:27)
        at org.apache.spark.rdd.RDD.computeOrReadCheckpoint(RDD.scala:306)
        at org.apache.spark.rdd.RDD.iterator(RDD.scala:270)
        at org.apache.spark.scheduler.ResultTask.runTask(ResultTask.scala:66)
        at org.apache.spark.scheduler.Task.run(Task.scala:89)
        at org.apache.spark.executor.Executor$TaskRunner.run(Executor.scala:227)
        at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:114
5)
        at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:61
5)
        at java.lang.Thread.run(Thread.java:745)
```

答：（原因是没装 X11 等依赖）

可以使用 iServer Support 中提供的一键化工具（810+版本），复制 support 目录到系统中。

```
./dependencies_check_and_install.sh install -y
```

5.9 iObjects for Spark 程序使用集群模式运行不成功，报错

```
java.lang.UnsatisfiedLinkError:
com.supermap.data.EnvironmentNative.jni_GetBasePath()Ljava/lang/String;
```

```
tu): java.lang.UnsatisfiedLinkError: com.supermap.data.EnvironmentNative.jni_GetBasePath
()Ljava/lang/String;
        at com.supermap.data.EnvironmentNative.jni_GetBasePath(Native Method)
        at com.supermap.data.Environment.getUGOBasePath(Environment.java:531)
        at com.supermap.geotools.LocalLogger.initConfig(LocalLogger.java:72)
        at com.supermap.geotools.LocalLogger.<clinit>(LocalLogger.java:14)
        at com.supermap.geotools.Convertor.<clinit>(Convertor.java:53)
        at com.supermap.bsp.rdd.BufferFeatureRDD.compute(BufferFeatureRDD.scala:27)
        at org.apache.spark.rdd.RDD.computeOrReadCheckpoint(RDD.scala:306)
        at org.apache.spark.rdd.RDD.iterator(RDD.scala:270)
        at org.apache.spark.scheduler.ResultTask.runTask(ResultTask.scala:66)
        at org.apache.spark.scheduler.Task.run(Task.scala:89)
        at org.apache.spark.executor.Executor$TaskRunner.run(Executor.scala:227)
        at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:114
5)
        at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:61
5)
        at java.lang.Thread.run(Thread.java:745)
```

答：使用 iObjects Java 中 libmawt.so 替换 jdk1.7.0_80/jre/lib/amd64/headless 中的 libmawt.so